


Data Administration in Information Systems

Database tuning (conclusion)

Database Tuning

- Second part for the course will address different tuning aspects, building on previous knowledge
 - Schema tuning
 - Query tuning
 - Index tuning
 - Lock and log tuning
 - Hardware and OS tuning
 - Database monitoring

Chapter 2



Tuning Considerations

- OS
 - Threads
 - Thread switching
 - Priorities
 - Virtual memory
 - Database buffer size
 - File system
 - Disk layout and access
- Hardware
 - Storage subsystem
 - Configuring the disk array
 - Using the controller cache
 - Hardware configuration

Threads


- A database system is an OS process with multiple threads
 - Threads to manage **database connections**
 - possibly using multiple network protocols
 - Threads to perform **system tasks**
 - e.g. checkpoints, running in the background
 - A pool of **worker threads**
 - to handle user requests, namely database queries

Threads (Cont.)

- Initiating (scheduling) a new thread takes some time
 - Pool of worker threads created at startup
- Switching between threads takes some time
 - Time slice of a thread should be long (e.g. 1 second) to dilute this cost
- Threads have priorities
 - Database system threads should not run at low priority

Threads (Cont.)

- Giving higher/lower priority to some transactions can backfire
 - may cause **priority inversion**
 - example
 - T_1 (higher priority) waiting for lock from T_3 (lower priority)
 - T_2 keeps running while T_3 (and therefore T_1) keep waiting



Threads (Cont.)

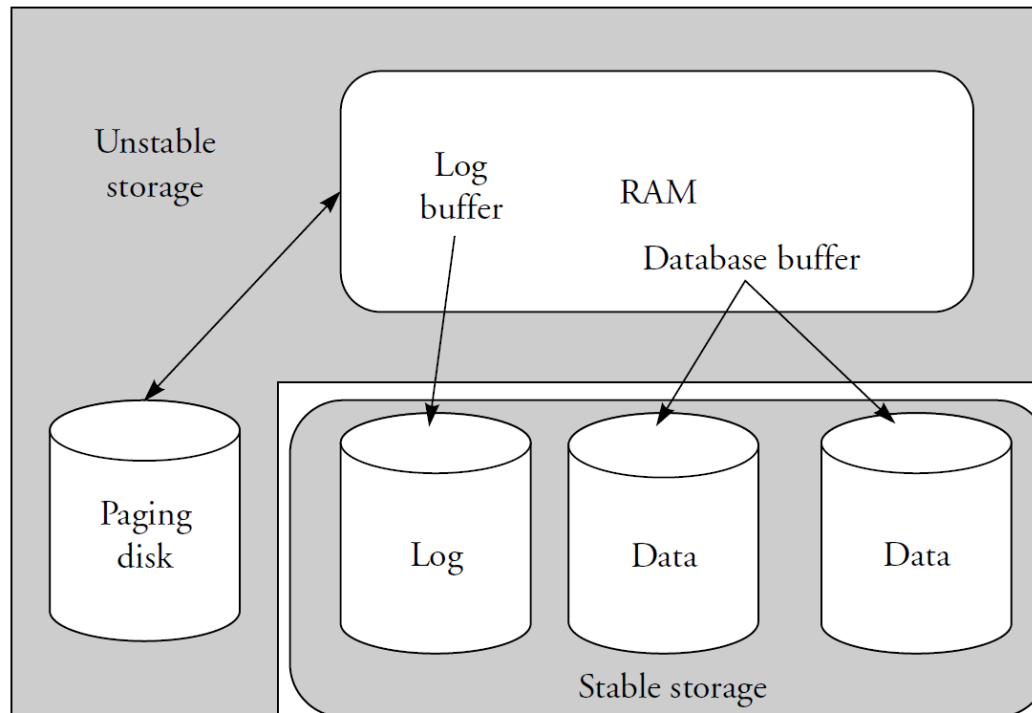
- Some systems use **priority inheritance**
 - When T_1 requests lock on X , the priority of T_3 is increased to that of T_1
 - If priority inheritance is unavailable, give same priority to all threads
- In SQL Server, there was **priority boosting**
 - Giving more priority to the database system than to the OS
 - Did not work out as expected, has been discontinued

Database Buffer

- Ideally, store as much as possible in RAM to avoid disk accesses
- Goal of memory tuning
 - Frequently read pages should rarely require disk accesses
- Logical vs. physical reads
 - Logical reads: pages that need to be accessed
 - Physical reads: pages that need to be retrieved from disk
- Logical vs. physical writes
 - Logical writes: changes to pages in the buffer (dirty pages)
 - Physical writes: dirty pages are written to disk
- Page replacements
 - Dirty pages are written to disk to free buffer space
 - New pages are read from disk into the free buffer space

Database Buffer (Cont.)

- When RAM becomes full, the OS may start paging to disk (**swap**)
 - OS paging is highly detrimental and should be avoided
 - Even worse is when the swap file is on the same disk
 - sequential access will be disrupted



Database Buffer (Cont.)


- Assuming OS paging and page replacements are low

$$\text{Hit ratio} = \frac{(\text{number of logical accesses}) - (\text{number of physical accesses})}{(\text{number of logical accesses})}$$

- Aim for hit ratio > 90%
- Run typical workload and check hit ratio
- Increase buffer size until hit ratio flattens out

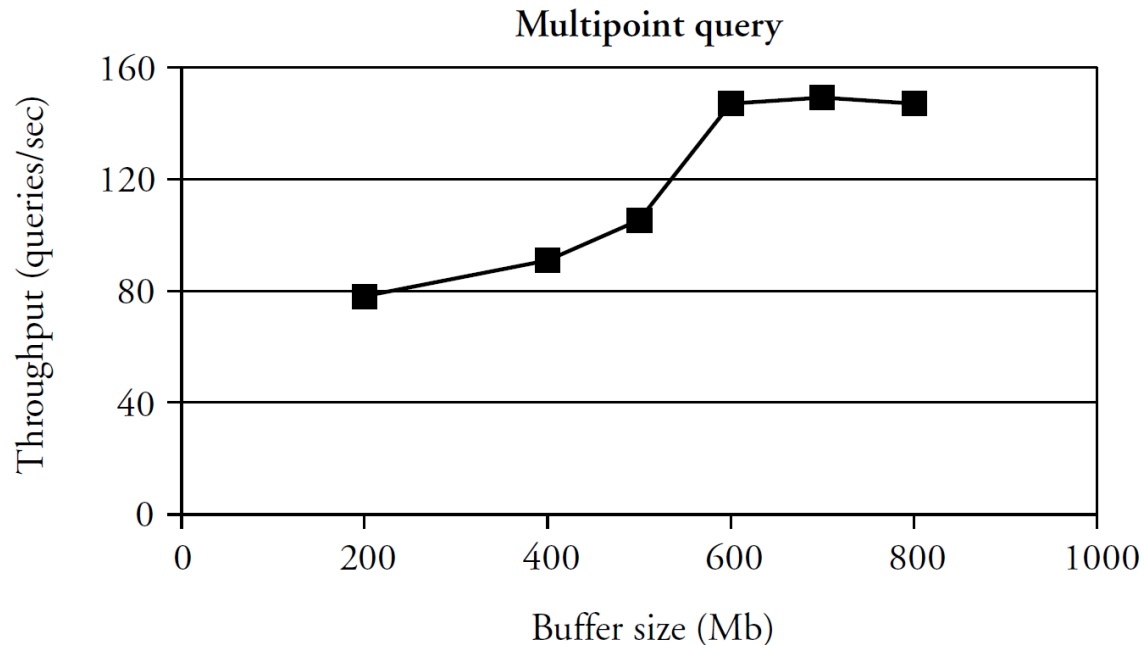
Database Buffer (Cont.)

- Experiment with full table scan
 - If buffer size is small, table must be read entirely from disk
 - LRU replacement strategy keeps evicting pages from memory
 - If buffer size is large, the entire table fits in memory
 - Query is processed entirely from RAM



Database Buffer (Cont.)

- Experiment with multipoint query
 - As the buffer size increases, more pages can be held in cache
 - LRU replacement strategy still evicts some pages from memory
 - After a certain point, all the desired pages fit in memory
 - Query is processed from RAM, no further improvement



File System


- Size of disk chunks allocated at one time
 - Some file systems call these **extents**
 - In SQL Server, an extent is 8 physically contiguous pages
 - Allocate large extents to files that need to be scanned often
 - Sequential files such as log or history also benefit from large extents

File System (Cont.)

- Usage factor of disk pages
 - Percentage of a page that is utilized, leaving some room for insertions
 - Depends on scan/insert ratio
 - High usage factor (90% or higher) is good for table scans
 - Low usage factor (70% or less) is good for frequent insertions
 - In SQL Server, there is also an **index fill factor**
 - percentage of space on each leaf-level page to be filled with data
 - we often rebuild indexes to reduce their internal fragmentation


File System (Cont.)

- Experiment with full table scan
 - Cold buffer to read entire table from disk
 - Throughput improves by about 10% when usage factor is increased from 70% to 100%



File System (Cont.)

- Pre-fetching pages from disk
 - Speed up table/index scans by reading ahead more pages from disk
 - Experiment with full table scan
 - Throughput improves by about 10% when the prefetching size is increased from 32 KB to 128 KB




Tuning Considerations


- OS
 - Threads
 - Thread switching
 - Priorities
 - Virtual memory
 - Database buffer size
 - File system
 - Disk layout and access
- Hardware
 - Storage subsystem
 - Configuring the disk array
 - Using the controller cache
 - Hardware configuration

RAID Levels


- RAID level 0
 - Block striping, non-redundant
- RAID level 1
 - Mirrored disks with block striping*
- RAID level 5
 - Block-interleaved distributed parity



(a) RAID 0: nonredundant striping



(b) RAID 1: mirrored disks



(c) RAID 5: block-interleaved distributed parity



* Some authors refer to mirrored disks *without* block striping as RAID level 1 and mirrored disks *with* block striping as RAID level 10.

RAID Levels (Cont.)

- Database log
 - RAID 1 is appropriate for log file(s)
 - Mirroring provides fault tolerance with high write throughput
 - Writes are synchronous and sequential; no benefit from striping
- Temporary Files
 - RAID 0 is appropriate for temporary tables or sorting files
 - No fault tolerance, high throughput; system can tolerate data loss
- Data and index files
 - RAID 5 provides fault tolerance and is best suited for read intensive apps

RAID Levels (Cont.)

- Experiments
 - Read performance
 - RAID 0, RAID 1, RAID 5 improve read performance with multiple disks
 - Write performance
 - Negative impact of RAID 5 for computing and writing parity block (hidden by controller cache)




Using the Controller Cache

- Disk controllers have memory that serves as cache
 - On read operations, the cache can be used for **read-ahead**
 - Similar to pre-fetching, but done internally in the hard drive
 - On write operations, the cache can be used for **write-back**
 - Write requests conclude as soon as data is written to the cache
 - In contrast with **write-through** mode, when it is written to disk

Using the Controller Cache (Cont.)

- Experiments with write-back mode
 - *Cache friendly*: the data volume is slightly larger than cache
 - *Cache unfriendly*: the data volume is 10x larger than cache
 - when cache is full, requests are serialized and waiting time increases
 - depends on disk access time and on length of waiting queue



Hardware Configuration

- Add memory
 - Allows increase in buffer size
 - Reduces load on disks
 - Increases the hit ratio
 - Reduces page replacement and OS paging

Hardware Configuration (Cont.)

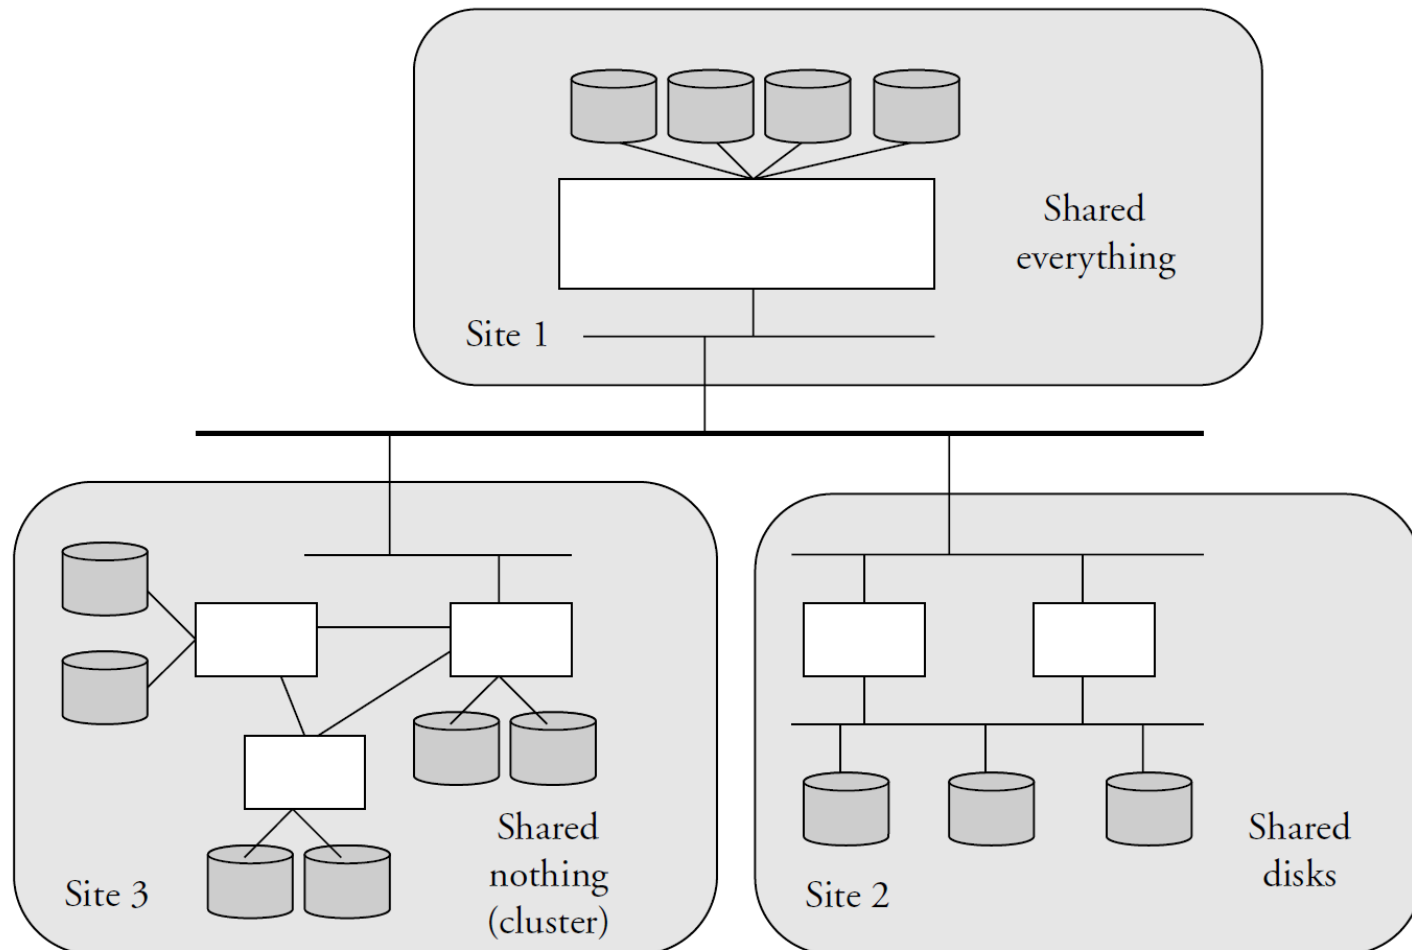
- Add disks
 - Put the log on a separate disk to ensure that writes are sequential
 - Use a different RAID level to achieve better write performance
 - e.g. switch from RAID 5 to RAID 1 for write-intensive apps
 - Partition large tables across several disks
 - Write-intensive apps should have non-clustered indexes to separate disk, because each modification updates those indexes
 - Read-intensive apps should partition tables across multiple disks to balance the read load

Hardware Configuration (Cont.)

- Add processors
 - Off-load non-database tasks to other processors
 - Provide computing power for data mining apps on copy of database
 - Connect many independent systems together by a high-speed network
 - Different options for sharing resources (memory, disks, processors)

Hardware Configuration (Cont.)


- Shared-everything, shared-disks, or shared-nothing environment



Database Tuning

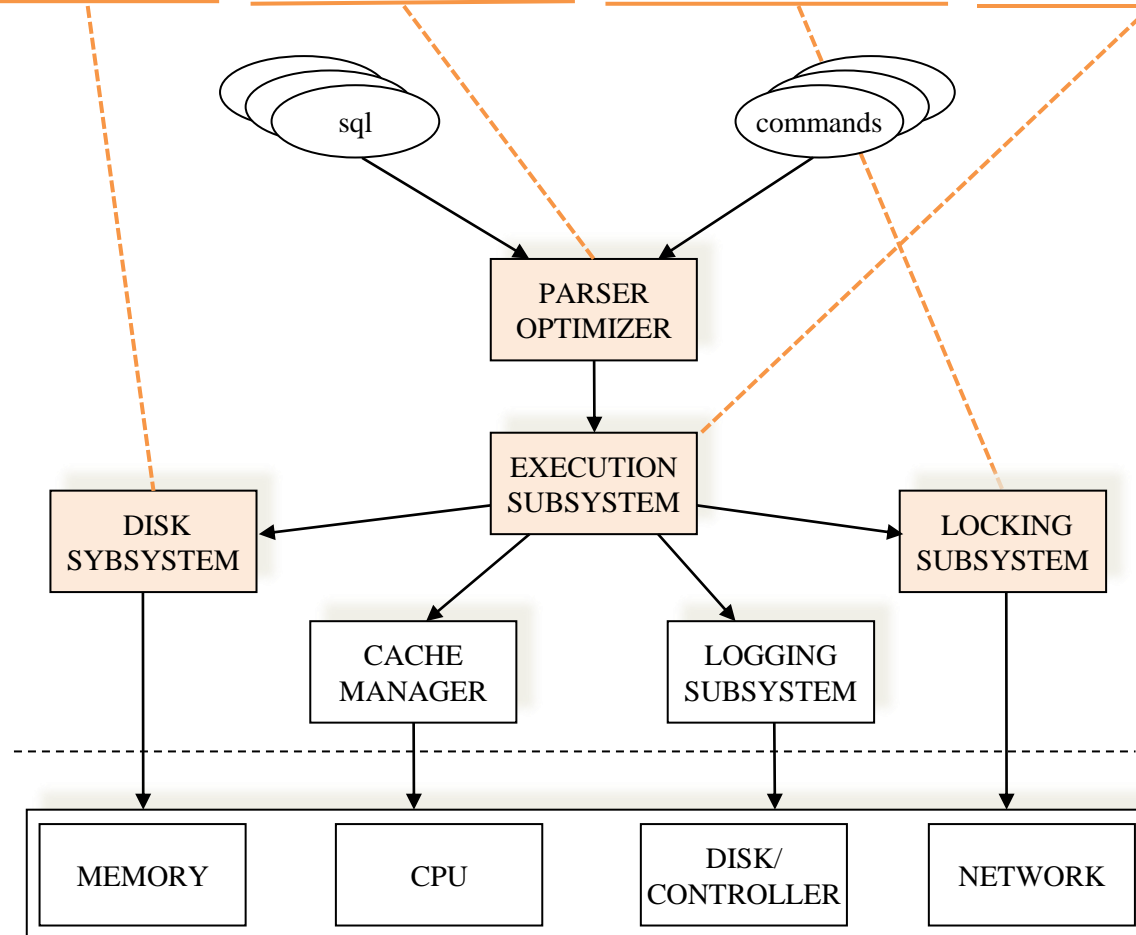
- Second part for the course will address different tuning aspects, building on previous knowledge
 - Schema tuning
 - Query tuning
 - Index tuning
 - Lock and log tuning
 - Hardware and OS tuning
 - Database monitoring

Chapter 7

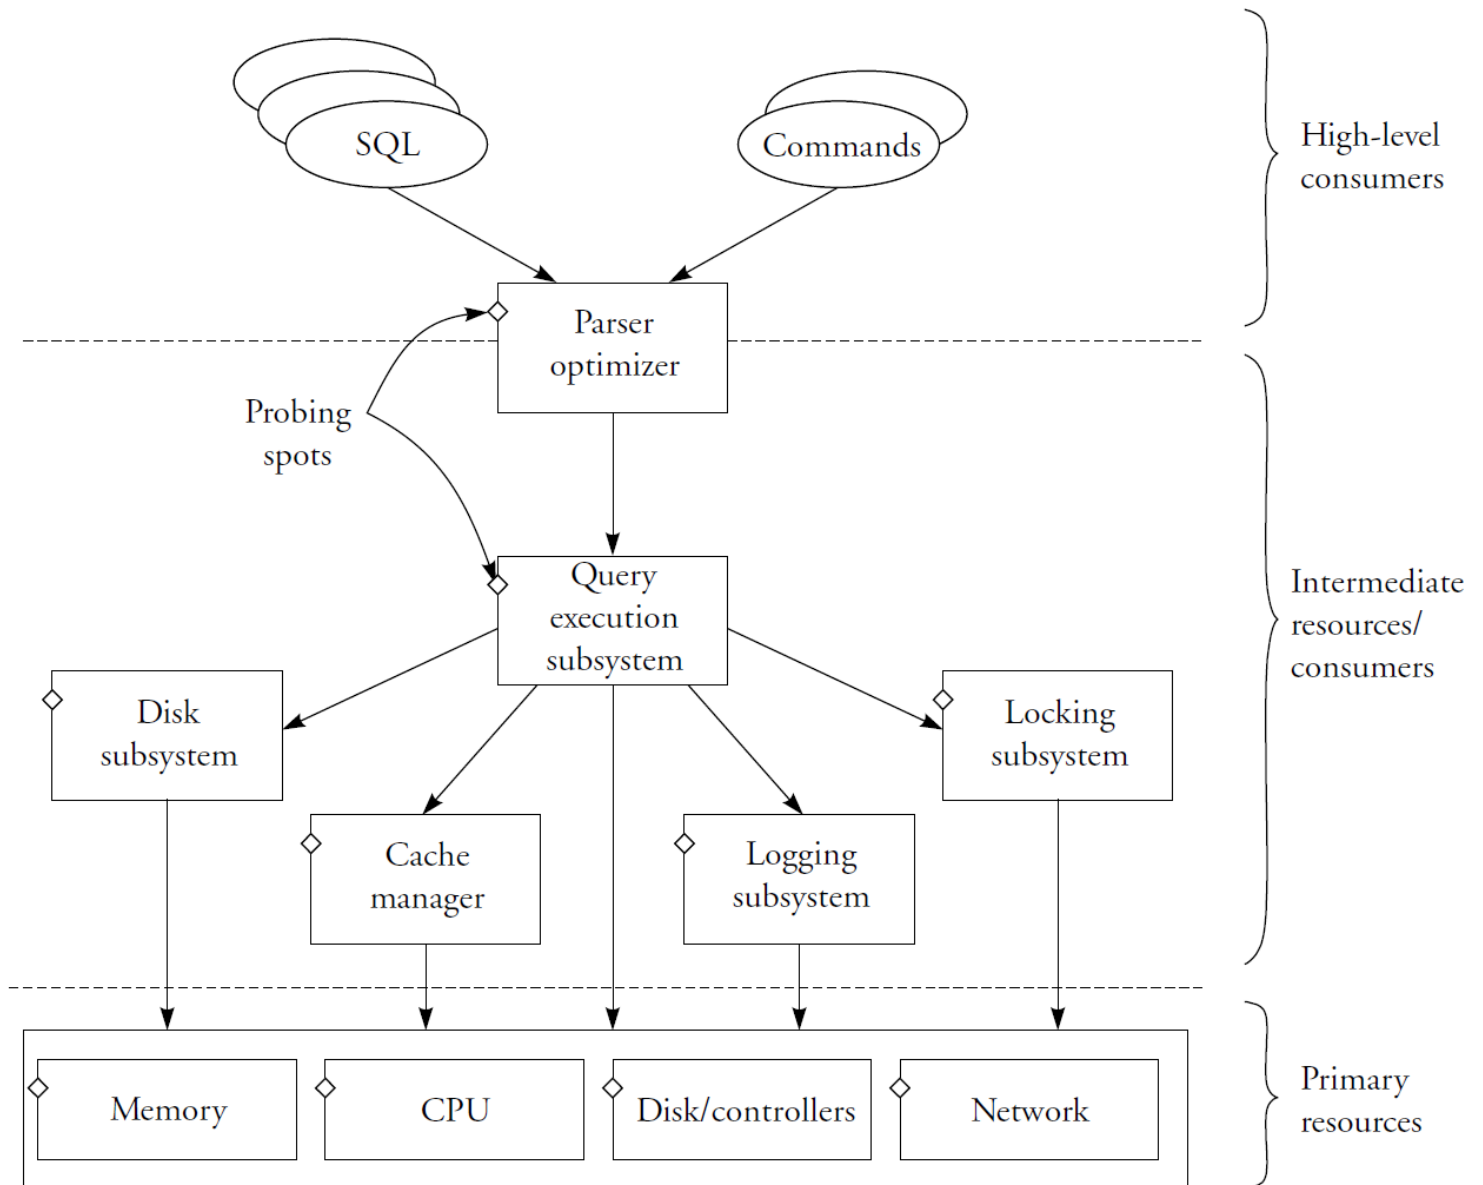


Database Monitoring

- Monitoring database performance requires **indicators**
 - These indicators can be collected at different points in the system
 - Disk accesses, query plan cost, number of locks, execution time, etc.



Consumption Chain

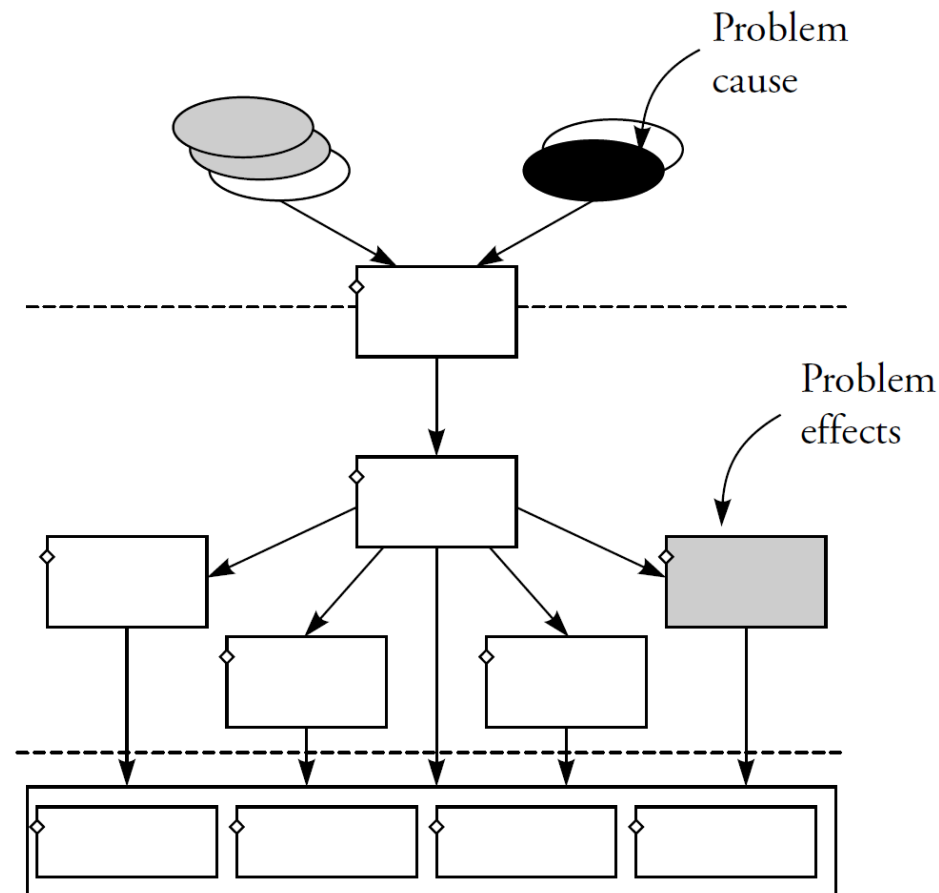


Consumption Chain (Cont.)

- Producer-consumer hierarchy
 - High-level consumers
 - Users or applications issuing SQL queries or database commands
 - Intermediate consumers/resources
 - Database subsystems that interact with each other those requests
 - Primary resources
 - Raw resources of the machine being managed by the OS

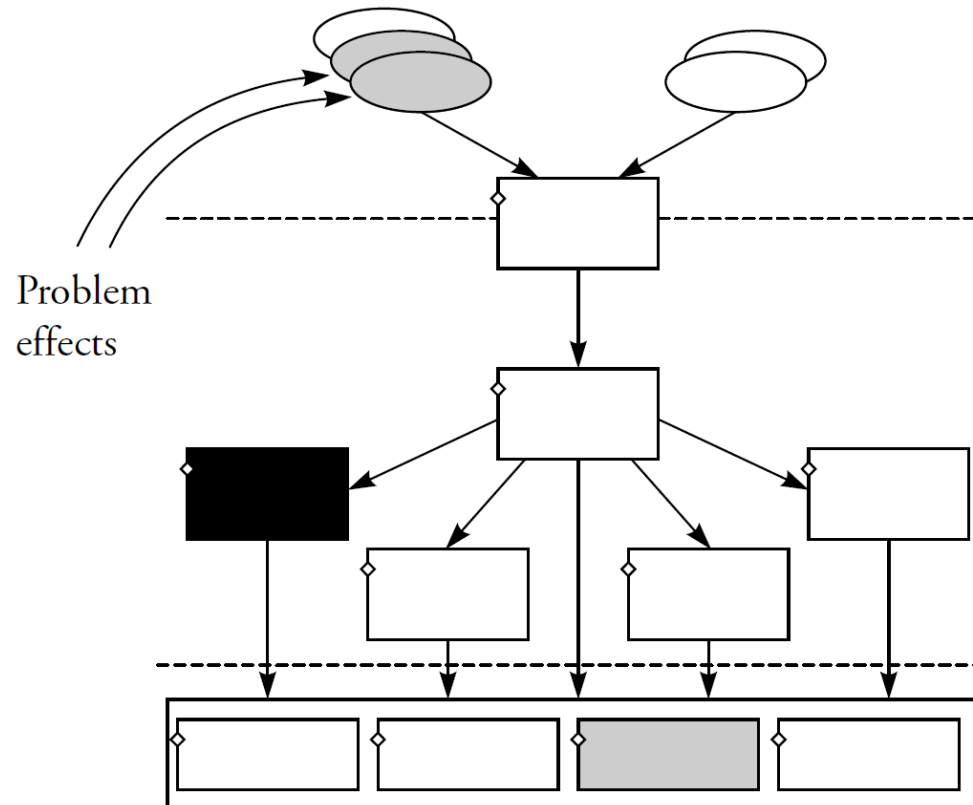
Performance problems

- Example 1
 - A high-level consumer monopolizing an intermediary resource
 - e.g. a query that locks an enormous number of rows




Performance problems (Cont.)

- Example 2
 - A poorly configured subsystem that exhausts a raw resource
 - e.g. a disk subsystem that stores everything on a single disk



Performance problems (Cont.)

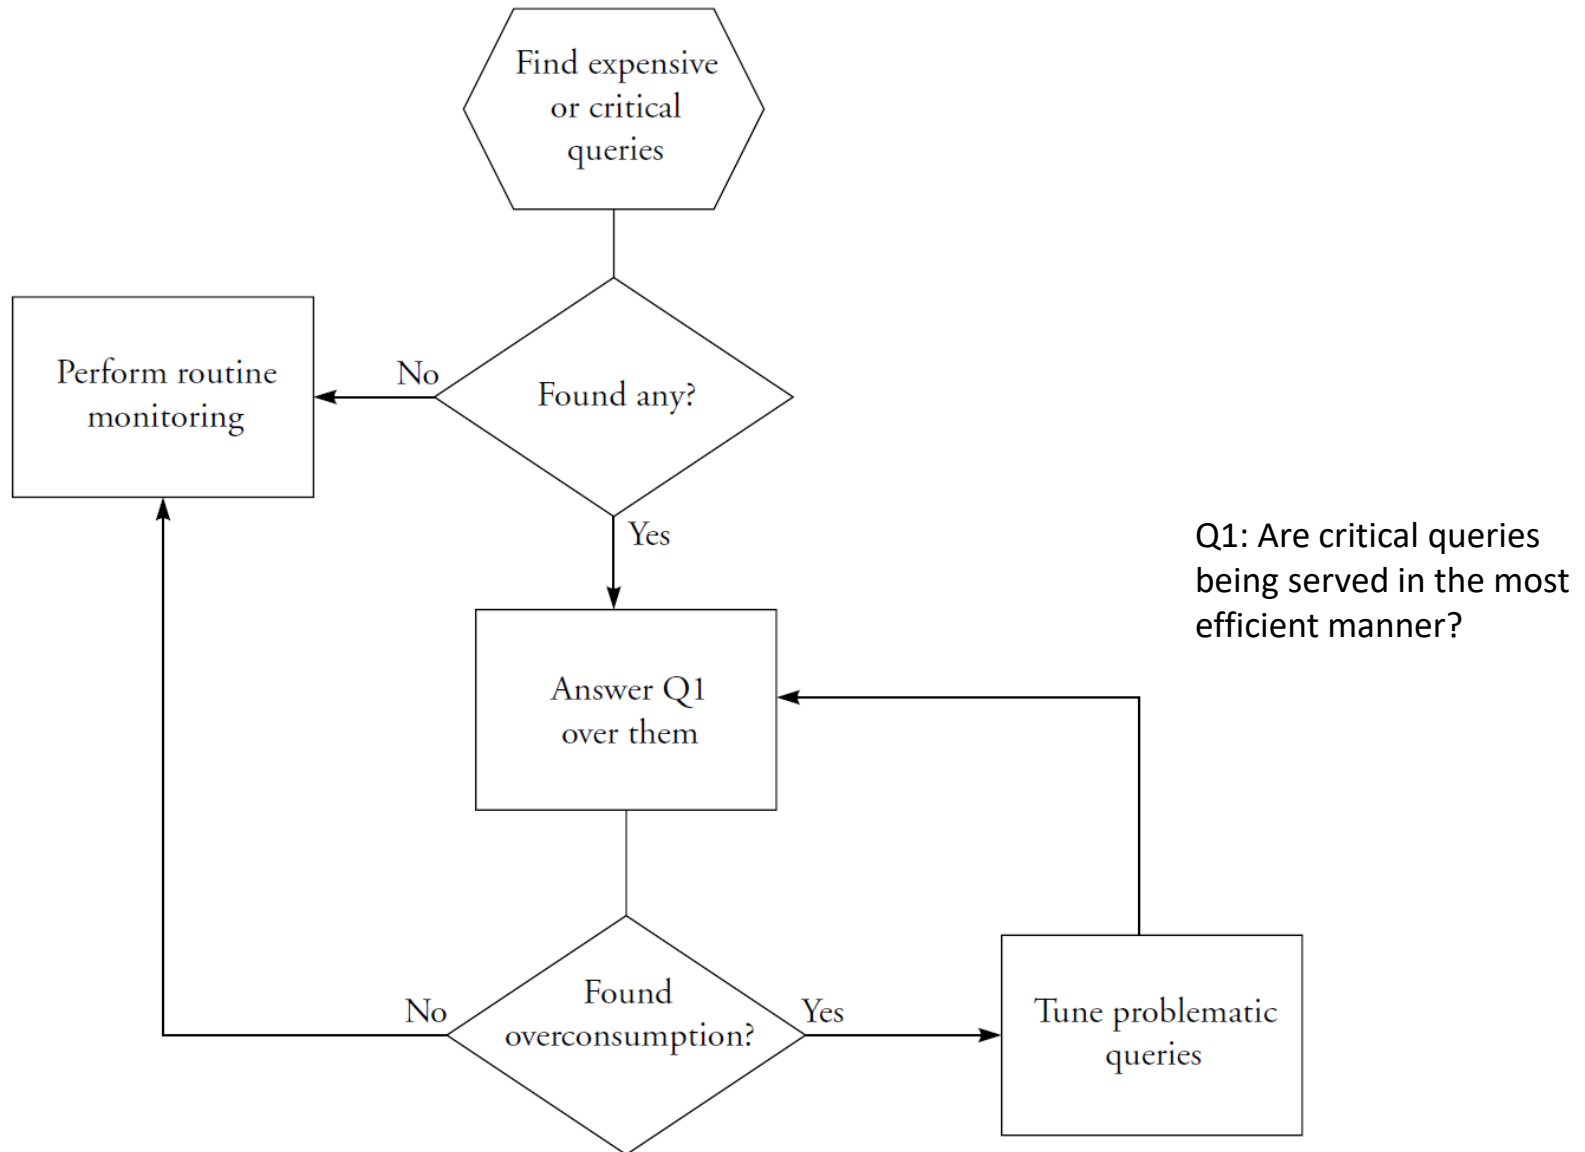
- Example 3
 - An overloaded primary resource can slow down an entire system
 - e.g. a CPU that is overly busy with non-database processes



Systematic Approach: Three Questions

1. Are critical queries being served in the most efficient manner?
 - High-level consumer question
2. Are database subsystems making optimal use of resources?
 - Intermediate consumer/resource question
3. Are there enough primary resources for the expected workload?
 - Primary resources question

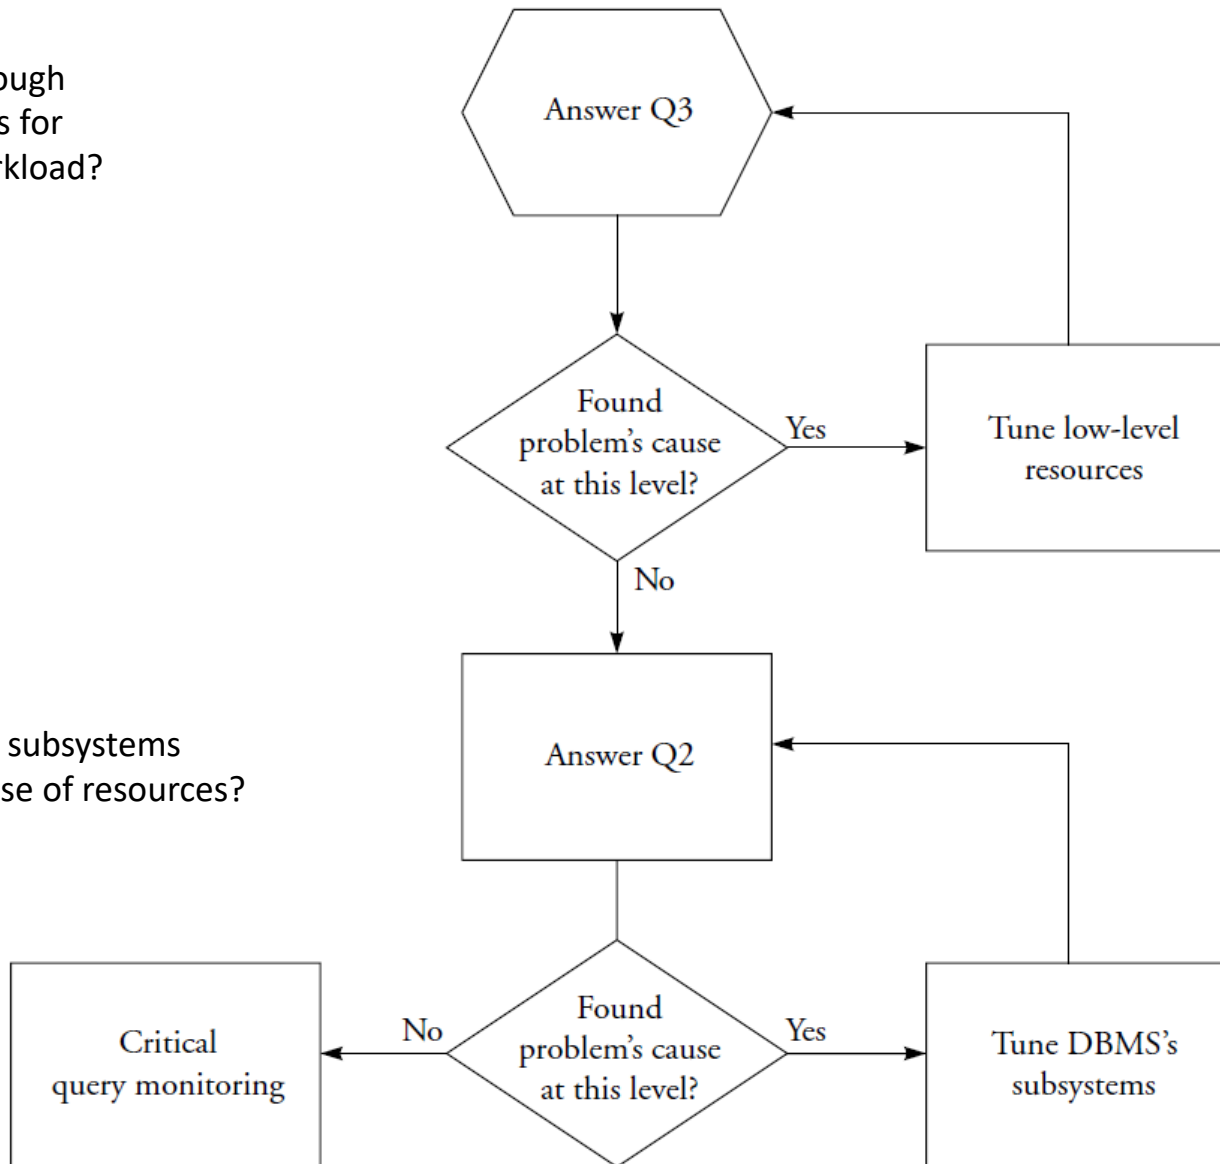
Critical Query Monitoring



Routine Monitoring

Q3: Are there enough primary resources for the expected workload?

Q2: Are database subsystems making optimal use of resources?



Monitoring Tools

- Query Plan Explainers
 - Shows the execution plan and estimated costs
- Performance Monitors
 - Tools that access the database internal counters and metrics
- Event Monitors
 - Record performance measures only when an event occurs

Monitoring Tools (Cont.)

- Query Plan Explainers
 - Shows the execution plan and estimated costs

The screenshot displays the Microsoft SQL Server Management Studio (SSMS) interface. The main window shows a query in the 'workload.sql' file, which is connected to the 'AdventureWorks2019' database. The query is as follows:

```
HAVING COUNT(*) > 1500;  
GO  
  
USE AdventureWorks2019;  
GO  
  
SELECT pp.FirstName, pp.LastName, e.NationalIDNumber  
FROM HumanResources.Employee AS e WITH (INDEX(AK_Employee_NationalIDNumber))  
JOIN Person.Person AS pp ON e.BusinessEntityID = pp.BusinessEntityID  
WHERE LastName = 'Johnson';  
GO  
  
USE AdventureWorks2019;  
GO  
  
SELECT ProductID, Total = SUM(LineTotal)  
FROM Sales.SalesOrderDetail  
GROUP BY ProductID  
HAVING SUM(LineTotal) > $1000000.00;
```

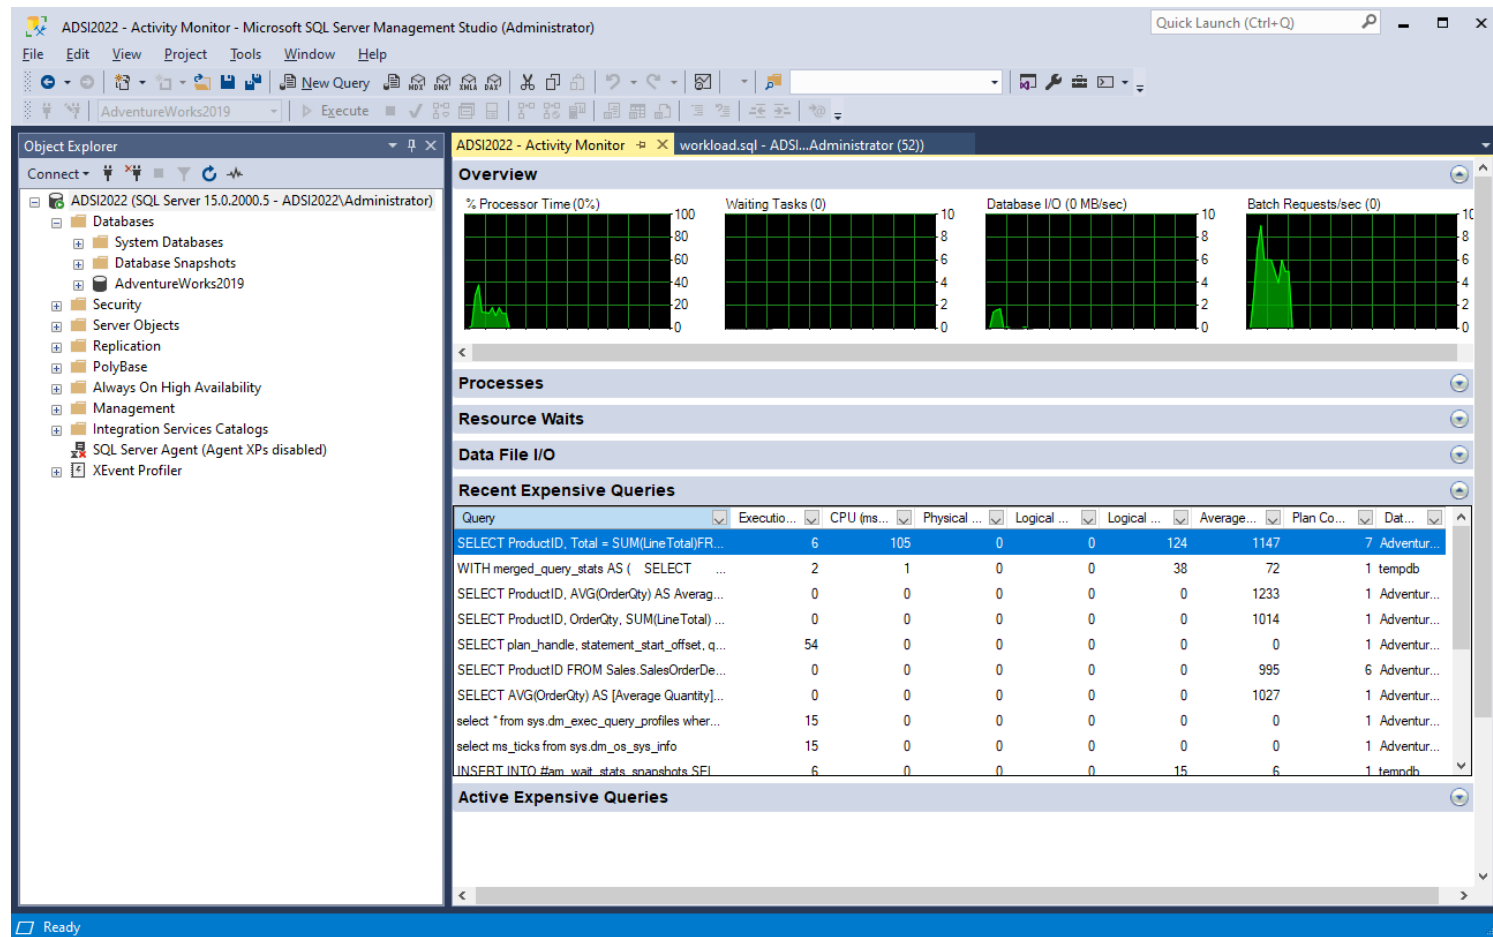
The 'Execution plan' tab is selected, showing the query cost (relative to the batch) as 100%. The execution plan diagram illustrates the following operations:

- SELECT** (Cost: 0 %): The root operation, which is a Hash Match (Inner Join).
- Hash Match (Inner Join)** (Cost: 76 %): This operation joins the results of the Index Seek and Index Scan.
- Index Seek (NonClustered)** (Cost: 11 %): Seeks for the [Person].[IX_Person_LastName_FirstName] index.
- Index Scan (NonClustered)** (Cost: 14 %): Scans the [Employee].[AK_Employee_NationalIDNumber] index.

The status bar at the bottom indicates that the query executed successfully, returning 3 rows in 00:00:00 seconds.

Monitoring Tools (Cont.)

- Performance Monitors
 - Tools that access the database internal counters and metrics



Monitoring Tools (Cont.)

- Event Monitors
 - Record performance measures only when an event occurs

| EventClass | TextData | ApplicationName | NTUserName | LoginName | CPU | Reads | Writes | Duration | ClientProcessID | SPID | StartTime |
|--------------------|---|-----------------|------------|-----------|-----|-------|--------|----------|-----------------|------|--------------|
| SQL:BatchCompleted | SELECT pp.FirstName, pp.LastName, e... | Microsoft SQ... | Adminis... | ADSI20... | 15 | 45 | 0 | 26 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | 0 | 0 | 0 | 0 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | SELECT ProductID, Total = SUM(LineTo... | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | SELECT ProductID, Total = SUM(LineTo... | Microsoft SQ... | Adminis... | ADSI20... | 360 | 1248 | 0 | 360 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | 0 | 0 | 0 | 0 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | SELECT pp.LastName, pp.FirstName, e... | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | SELECT pp.LastName, pp.FirstName, e... | Microsoft SQ... | Adminis... | ADSI20... | 15 | 50 | 0 | 8 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | 0 | 0 | 0 | 0 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | SELECT ProductID, OrderQty, SUM(Line... | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | SELECT ProductID, OrderQty, SUM(Line... | Microsoft SQ... | Adminis... | ADSI20... | 110 | 1280 | 0 | 118 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | 0 | 0 | 0 | 0 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | SELECT ProductID, AVG(UnitPrice) AS ... | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | SELECT ProductID, AVG(UnitPrice) AS ... | Microsoft SQ... | Adminis... | ADSI20... | 31 | 1276 | 0 | 36 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | 0 | 0 | 0 | 0 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | SELECT ProductID, Total = SUM(LineTo... | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | SELECT ProductID, Total = SUM(LineTo... | Microsoft SQ... | Adminis... | ADSI20... | 359 | 1279 | 0 | 362 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | 0 | 0 | 0 | 0 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | SELECT BusinessEntityID, JobTitle, H... | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | SELECT BusinessEntityID, JobTitle, H... | Microsoft SQ... | Adminis... | ADSI20... | 16 | 48 | 0 | 32 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | 0 | 0 | 0 | 0 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | SELECT ProductID, Total = SUM(LineTo... | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | SELECT ProductID, Total = SUM(LineTo... | Microsoft SQ... | Adminis... | ADSI20... | 375 | 1279 | 0 | 388 | 7084 | 55 | 2022-04-1... |

```
SELECT ProductID, AVG(UnitPrice) AS [Average Price]
FROM Sales.SalesOrderDetail
WHERE OrderQty > 15
GROUP BY ProductID
ORDER BY AVG(UnitPrice);
```

Trace is running. Ln 175, Col 2 Rows: 187 Connections: 1

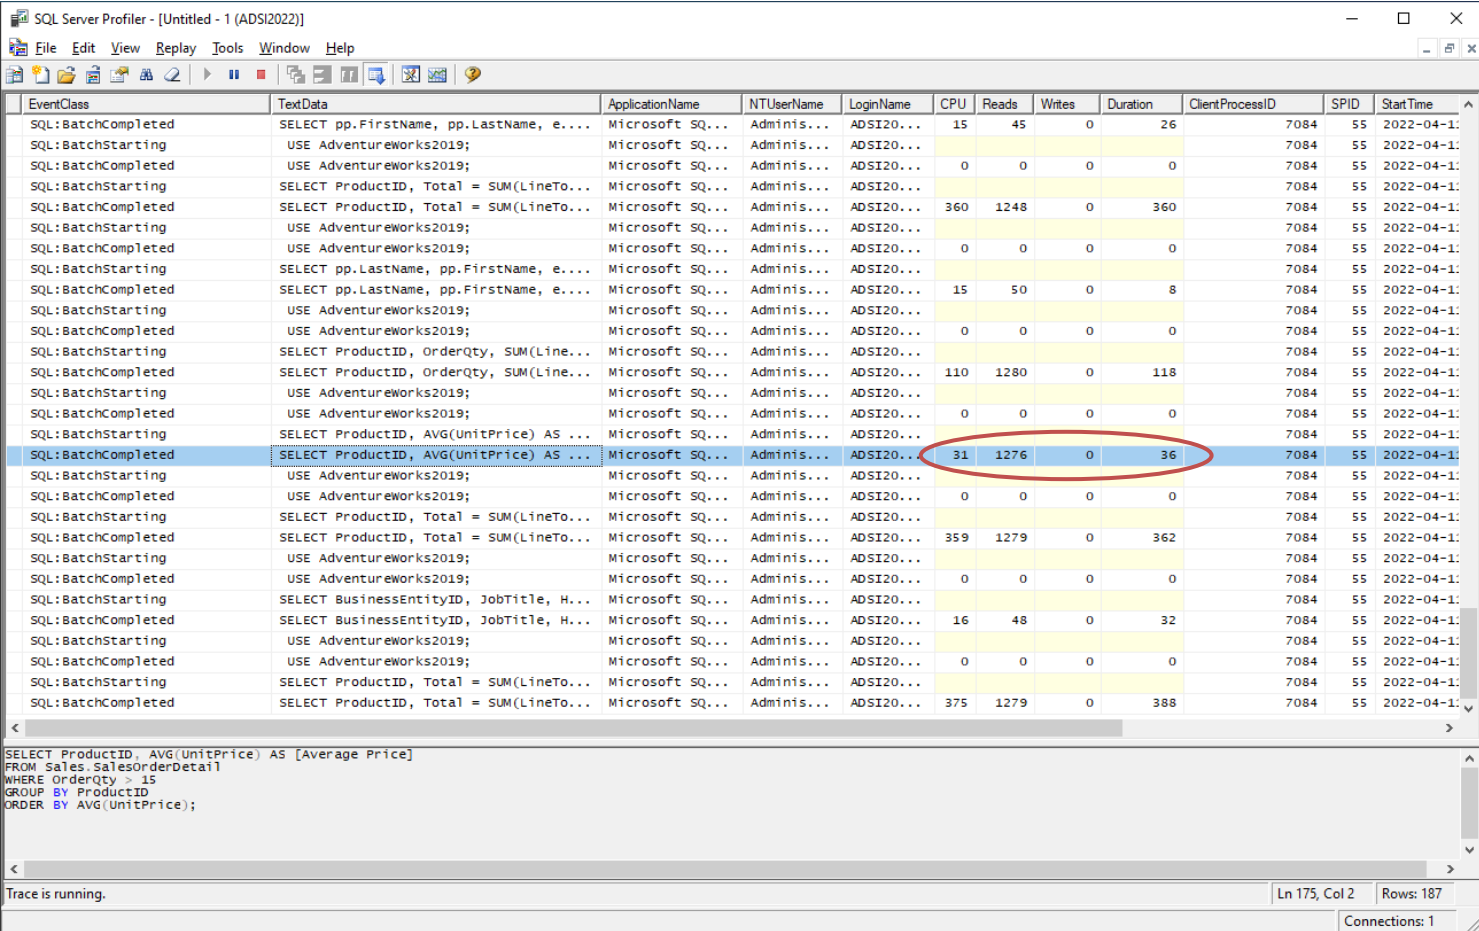
Investigating High-Level Consumers

- Answer question 1
 - Are critical queries being served in the most efficient manner?
1. Identify the critical queries
 2. Analyze the execution plan
 3. Profile the execution

Investigating High-Level Consumers (Cont.)

1. Identify the critical queries

- Use Event Monitor to find end-of-statement with execution measures



| EventClass | TextData | ApplicationName | NTUserName | LoginName | CPU | Reads | Writes | Duration | ClientProcessID | SPID | StartTime |
|--------------------|---|-----------------|------------|-----------|-----|-------|--------|----------|-----------------|------|--------------|
| SQL:BatchCompleted | SELECT pp.FirstName, pp.LastName, e... | Microsoft SQ... | Adminis... | ADSI20... | 15 | 45 | 0 | 26 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | 0 | 0 | 0 | 0 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | SELECT ProductID, Total = SUM(LineTo... | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | SELECT ProductID, Total = SUM(LineTo... | Microsoft SQ... | Adminis... | ADSI20... | 360 | 1248 | 0 | 360 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | 0 | 0 | 0 | 0 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | SELECT pp.LastName, pp.FirstName, e... | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | SELECT pp.LastName, pp.FirstName, e... | Microsoft SQ... | Adminis... | ADSI20... | 15 | 50 | 0 | 8 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | 0 | 0 | 0 | 0 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | SELECT ProductID, OrderQty, SUM(Line... | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | SELECT ProductID, OrderQty, SUM(Line... | Microsoft SQ... | Adminis... | ADSI20... | 110 | 1280 | 0 | 118 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | 0 | 0 | 0 | 0 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | SELECT ProductID, AVG(UnitPrice) AS ... | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | SELECT ProductID, AVG(UnitPrice) AS ... | Microsoft SQ... | Adminis... | ADSI20... | 31 | 1276 | 0 | 36 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | 0 | 0 | 0 | 0 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | SELECT ProductID, Total = SUM(LineTo... | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | SELECT ProductID, Total = SUM(LineTo... | Microsoft SQ... | Adminis... | ADSI20... | 359 | 1279 | 0 | 362 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | 0 | 0 | 0 | 0 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | SELECT BusinessEntityID, JobTitle, H... | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | SELECT BusinessEntityID, JobTitle, H... | Microsoft SQ... | Adminis... | ADSI20... | 16 | 48 | 0 | 32 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | USE Adventureworks2019; | Microsoft SQ... | Adminis... | ADSI20... | 0 | 0 | 0 | 0 | 7084 | 55 | 2022-04-1... |
| SQL:BatchStarting | SELECT ProductID, Total = SUM(LineTo... | Microsoft SQ... | Adminis... | ADSI20... | | | | | 7084 | 55 | 2022-04-1... |
| SQL:BatchCompleted | SELECT ProductID, Total = SUM(LineTo... | Microsoft SQ... | Adminis... | ADSI20... | 375 | 1279 | 0 | 388 | 7084 | 55 | 2022-04-1... |

```
SELECT ProductID, AVG(UnitPrice) AS [Average Price]
FROM Sales.SalesOrderDetail
WHERE OrderQty > 15
GROUP BY ProductID
ORDER BY AVG(UnitPrice);
```

Trace is running. Ln 175, Col 2 Rows: 187 Connections: 1

Investigating High-Level Consumers (Cont.)

2. Analyze the execution plan

- Use Query Plan Explorer to analyze the relative cost of each operation

The screenshot displays the Microsoft SQL Server Management Studio (SSMS) interface. The main window shows a query in the 'workload.sql' file, which is connected to the 'AdventureWorks2019' database. The query is as follows:

```
HAVING COUNT(*) > 1500;  
GO  
  
USE AdventureWorks2019;  
GO  
  
SELECT pp.FirstName, pp.LastName, e.NationalIDNumber  
FROM HumanResources.Employee AS e WITH (INDEX(AK_Employee_NationalIDNumber))  
JOIN Person.Person AS pp ON e.BusinessEntityID = pp.BusinessEntityID  
WHERE LastName = 'Johnson';  
GO  
  
USE AdventureWorks2019;  
GO  
  
SELECT ProductID, Total = SUM(LineTotal)  
FROM Sales.SalesOrderDetail  
GROUP BY ProductID  
HAVING SUM(LineTotal) > $100000.00;
```

The 'Execution plan' tab is selected, showing the query cost (relative to the batch): 100%. The execution plan diagram illustrates the following operations:

- SELECT**: Cost: 0 %
- Hash Match (Inner Join)**: Cost: 76 %
- Index Seek (NonClustered)**: [Person].[IX_Person_LastName_FirstN...], Cost: 11 %
- Index Scan (NonClustered)**: [Employee].[AK_Employee_NationalIDN...], Cost: 14 %

The status bar at the bottom indicates: 'Query executed successfully. ADSI2022 (15.0 RTM) ADSI2022\Administrator... AdventureWorks2019 00:00:00 3 rows'.

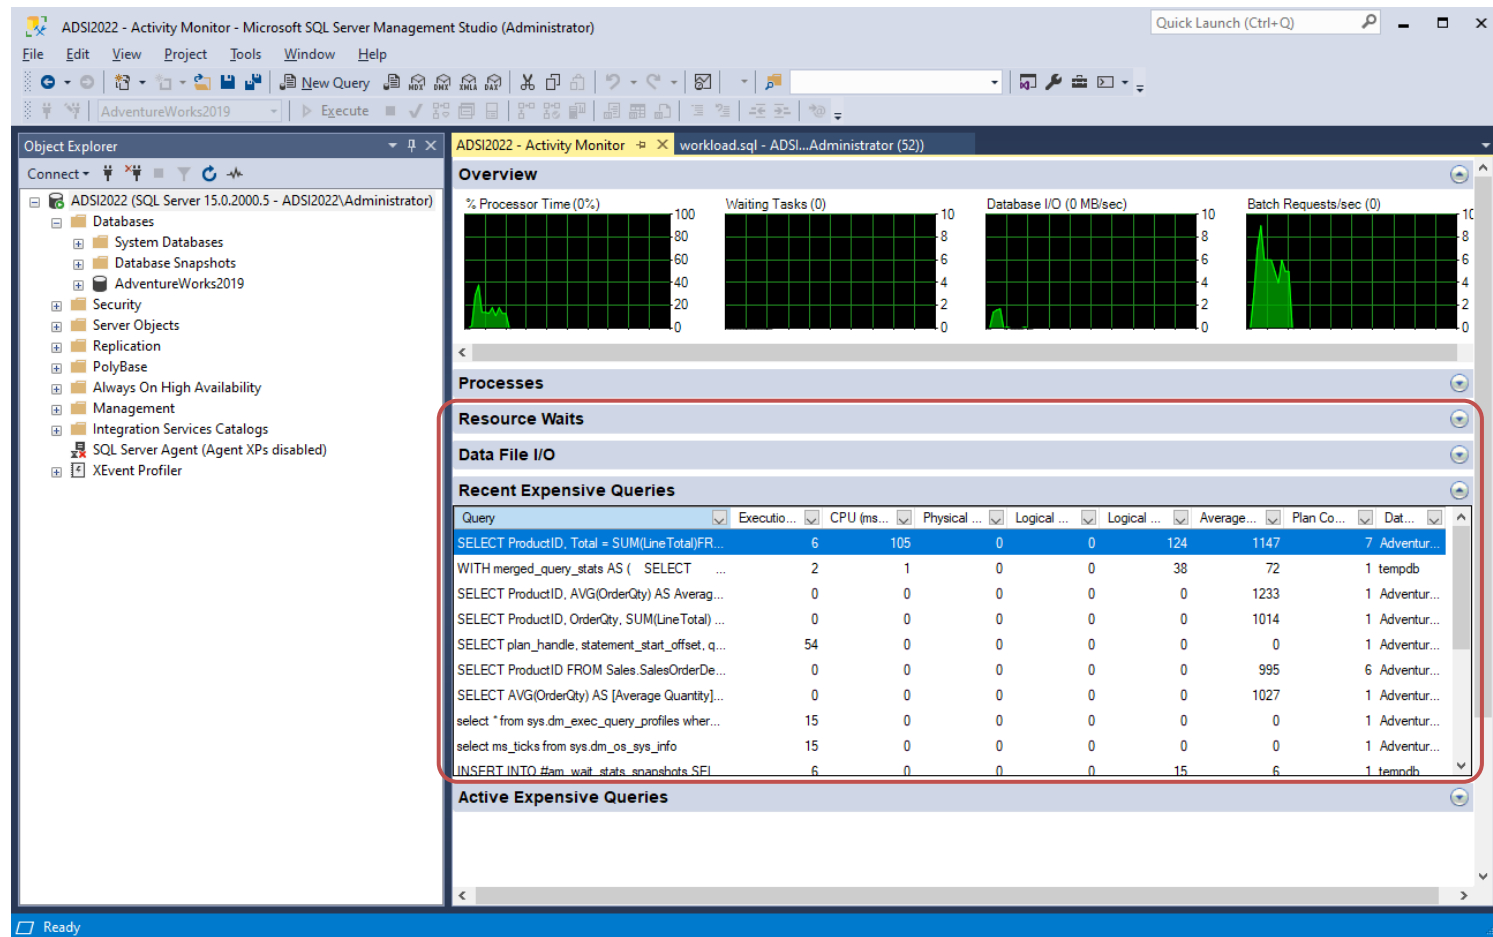
Investigating High-Level Consumers (Cont.)

- In the execution plan, pay attention to:
 - Access methods
 - sequential scan, index lookup, etc.
 - Sorts
 - caused by ORDER BY, GROUP BY or DISTINCT
 - Intermediate results
 - materialization to temporaries
 - Order of operations
 - joins, sorts, aggregations, filtering
 - Algorithms used in operations
 - types of join, etc.

Investigating High-Level Consumers (Cont.)

3. Profile the execution

- Use Performance Monitor to analyze duration and resource consumption



Investigating High-Level Consumers (Cont.)

- Duration involves 3 indicators:
 - Elapsed time
 - The time it took to process the query as perceived by a user
 - CPU time
 - The time that was actually used by the CPU to process the query
 - Wait time
 - The time the query was waiting for a resources to become available
- Resource consumption includes:
 - I/O
 - Physical and logical reads/writes
 - Locking
 - Number of locks, lock escalations, deadlocks/timeouts, total wait time
 - SQL activity
 - Number of sorts and temporary area usage

Investigating High-Level Consumers (Cont.)

- Two common scenarios
 - **Elapsed time close to CPU time**
 - Probably difficult to optimize any further
 - **Discrepancy between elapsed time and CPU time**
 - Points to a problem in resource consumption
 - Possibly a contention problem or a poorly performing resource
 - Run the query in isolation to investigate the cause

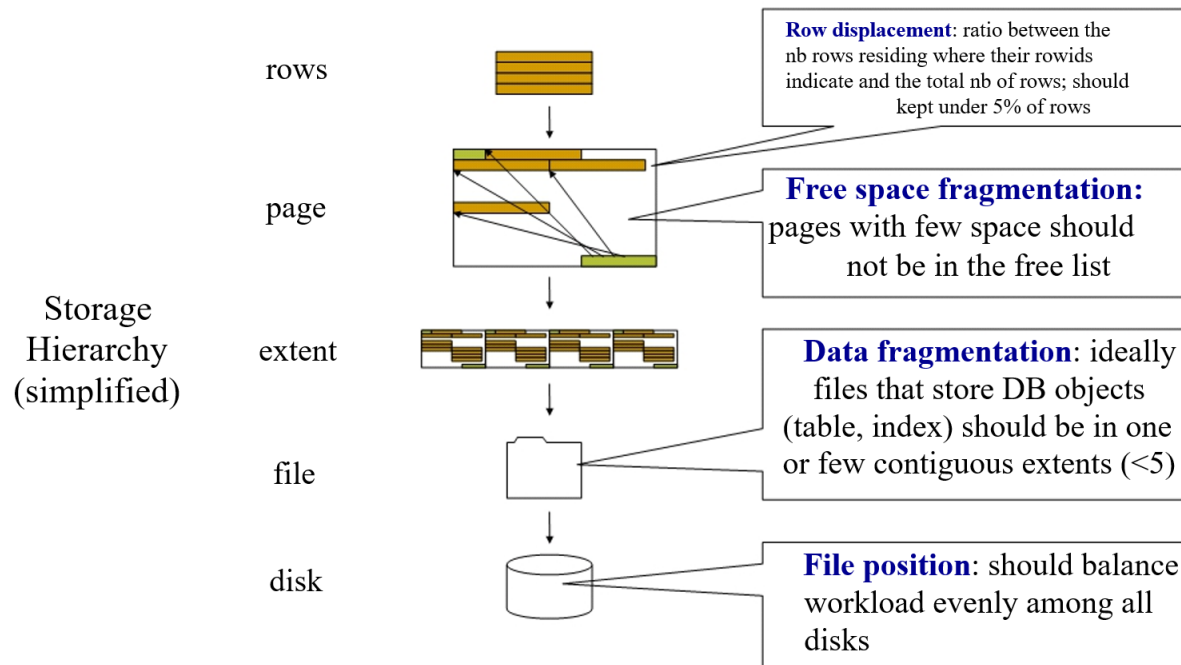
Investigating Intermediate Resources/Consumers

- Answer question 2
 - Are database subsystems making optimal use of resources?
1. Disk subsystem
 2. Buffer manager
 3. Locking subsystem
 4. Logging subsystem

Investigating Intermediate Resources/Consumers (Cont.)

1. Disk subsystem

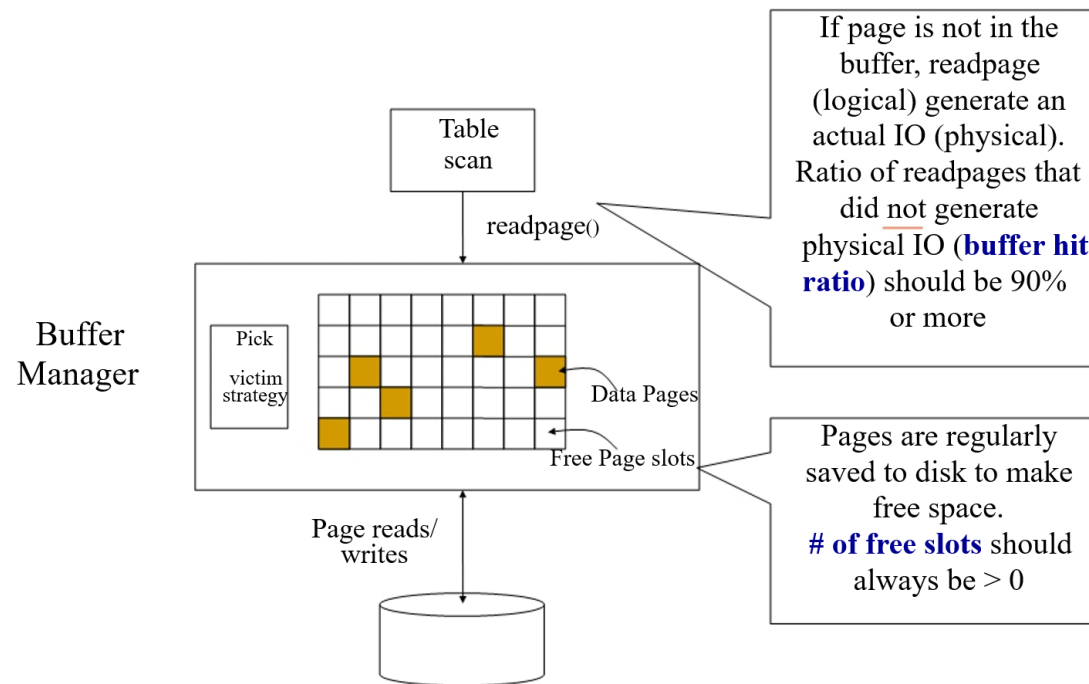
- A table should be stored contiguously in a physical disk
 - Avoid free space between records (**data fragmentation**)
- Table records should be stored in their correct order
 - Avoid records out of place (**row displacement**)
- Periodic file reorganization may be necessary



Investigating Intermediate Resources/Consumers (Cont.)

2. Buffer Manager

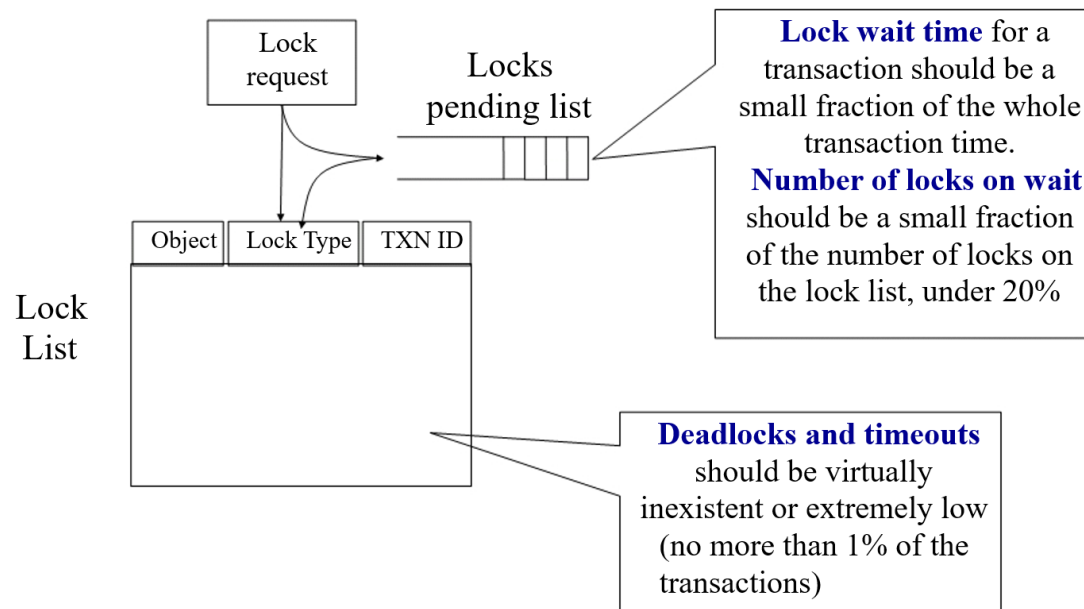
- Two main performance indicators to monitor
 - **Hit ratio** – percentage of times that requested page is already in buffer
 - **Number of free pages** – how much space is left in the buffer
- In SQL Server, these and other metrics can be obtained from system views



Investigating Intermediate Resources/Consumers (Cont.)

3. Locking subsystem

- Useful indicators
 - **Average lock wait time**
 - **Number of locks on wait**
 - **Number of deadlocks or timeouts**
- SQL Server provides comprehensive wait statistics through system views



~ ~ ~

Investigating Intermediate Resources/Consumers (Cont.)

4. Logging subsystem

- Useful indicators
 - **Number of log waits** – ensure log can keep up with transactions
 - **Log expansions** or **log archives** – due to lack of space
 - **Log cache hit ratio** – analogous to buffer cache hit ratio
- Log waits > 0 means transactions are being held due to log writes

Investigating Primary Resources

- Answer question 3
 - Are there enough primary resources for the expected workload?

1. CPU


2. Disks

3. Memory

Investigating Primary Resources (Cont.)

1. CPU

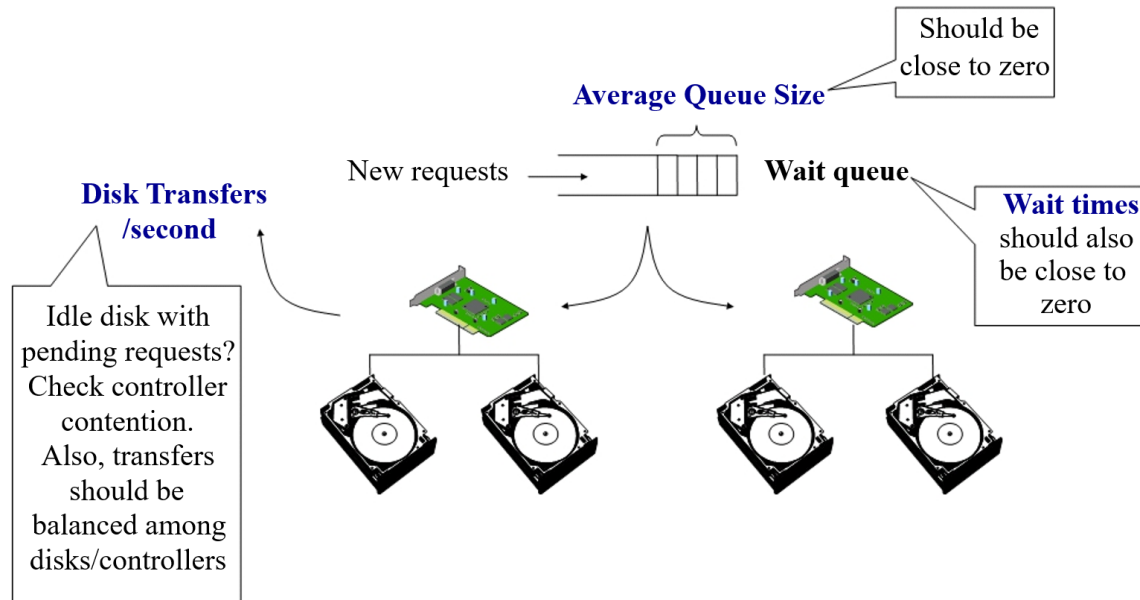
- Main indicator
 - **Percentage of utilization**
- Use OS task manager to monitor CPU utilization
- Identify whether processes are database or non-database related
- Check CPU utilization of system (OS) processes in idle state



Investigating Primary Resources (Cont.)

2. Disks

- Main indicators
 - **Average size of the waiting queue**
 - **Average time taken to service a request**
 - **Bytes transferred per second**
- Disk utilization can be monitored with OS utilities



Investigating Primary Resources (Cont.)

3. Memory

- Some indicators
 - **Number of page faults/time**
 - **Percentage of paging file in use**
- Size of paging/swap file is an indication of how much memory is lacking

